

```
In [1]: import pandas as pd

retail_sales_df=pd.read_csv('us_retail_sales.csv')
retail_sales_df.head()
```

```
Out[1]:
```

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT
0	1992	146925	147223	146805	148032	149010	149800	150761.0	151067.0	152588.0	153521.0
1	1993	157555	156266	154752	158979	160605	160127	162816.0	162506.0	163258.0	164685.0
2	1994	167518	169649	172766	173106	172329	174241	174781.0	177295.0	178787.0	180561.0
3	1995	182413	179488	181013	181686	183536	186081	185431.0	186806.0	187366.0	186565.0
4	1996	189135	192266	194029	194744	196205	196136	196187.0	196218.0	198859.0	200509.0

```
In [2]: retail_sales_df.shape
```

```
Out[2]: (30, 13)
```

```
In [3]: # Transpose Data
Transpose_retail_sales_df=pd.melt(retail_sales_df, id_vars=['YEAR'])
Transpose_retail_sales_df.head()
```

```
Out[3]:
```

	YEAR	variable	value
0	1992	JAN	146925.0
1	1993	JAN	157555.0
2	1994	JAN	167518.0
3	1995	JAN	182413.0
4	1996	JAN	189135.0

```
In [4]: # Check Nulls
Transpose_retail_sales_df.isnull().sum()
```

```
Out[4]: YEAR      0
variable    0
value      6
dtype: int64
```

```
In [5]: # Drop Nulls
Transpose_retail_sales_df.dropna(inplace=True)
Transpose_retail_sales_df.isnull().sum()
Transpose_retail_sales_df.head()
```

```
Out[5]:
```

	YEAR	variable	value
--	------	----------	-------

	YEAR	variable	value
0	1992	JAN	146925.0
1	1993	JAN	157555.0
2	1994	JAN	167518.0
3	1995	JAN	182413.0
4	1996	JAN	189135.0

In [6]:

```
import datetime as dt

#convert date and hour to str, concatenate them and then convert them to datetime format
Transpose_retail_sales_df['new_date'] = Transpose_retail_sales_df[['YEAR','variable']].
```

In [7]:

```
# Order the values and set the new date field as the index

final_retail_sales_df=Transpose_retail_sales_df[['new_date','value']]
final_retail_sales_df.sort_values(by=['new_date'], inplace=True)
final_retail_sales_df.reset_index(drop=True, inplace=True)
final_retail_sales_df.set_index('new_date',inplace=True)
final_retail_sales_df.head(100)
```

<ipython-input-7-e08b33fa83ee>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
final_retail_sales_df.sort_values(by=['new_date'], inplace=True)

Out[7]:


	value
new_date	
1992-01-01	146925.0
1992-02-01	147223.0
1992-03-01	146805.0
1992-04-01	148032.0
1992-05-01	149010.0
...	...
1999-12-01	245485.0
2000-01-01	243436.0
2000-02-01	247133.0
2000-03-01	249825.0
2000-04-01	245831.0

100 rows × 1 columns

```
In [8]: # Line plot

import matplotlib.pyplot as plt

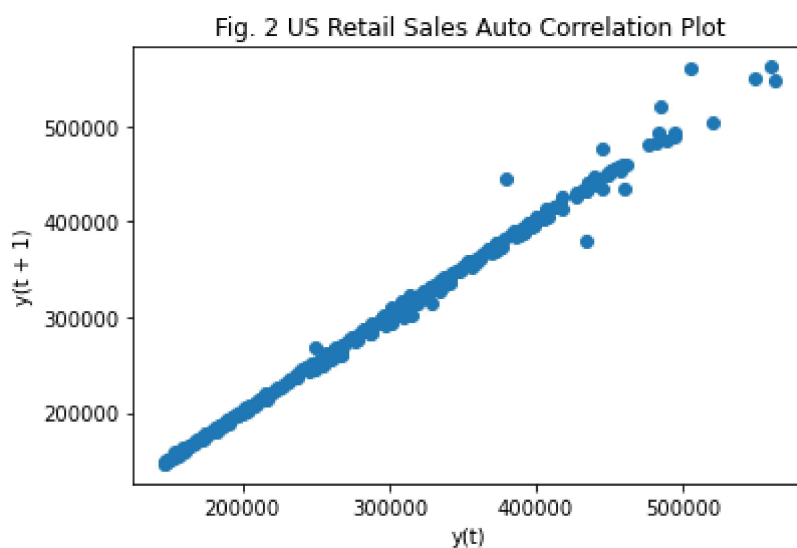
plt.plot(final_retail_sales_df)
plt.xlabel("Year")
plt.ylabel("Sales")
plt.title("Fig. 1 US Retail Sales")
plt.show()
```



From the line plot we can infer that sales grew strongly over time although there were noticeable dips during 2008 and 2020 when there was an economic slowdown

```
In [9]: #Check Auto Correlation

pd.plotting.lag_plot(final_retail_sales_df, lag=1)
plt.title("Fig. 2 US Retail Sales Auto Correlation Plot")
plt.show()
```



We can see a large ball of observations along a diagonal line of the plot. It clearly shows strong correlation.

```
In [10]:
```

```

from pandas import read_csv
from pandas import DataFrame
from pandas import concat

values = DataFrame(final_retail_sales_df)
dataframe = concat([values.shift(1), values], axis=1)
dataframe.columns = ['t-1', 't+1']
result = dataframe.corr()
print(result)

```

```

          t-1          t+1
t-1  1.000000  0.997354
t+1  0.997354  1.000000

```

In [11]: *# create and evaluate a static autoregressive model*

```

from statsmodels.tsa.ar_model import AutoReg
from sklearn.metrics import mean_squared_error
from math import sqrt

# split dataset
X = final_retail_sales_df.values
train, test = X[1:len(X)-12], X[len(X)-12:]

```

In [13]:

```

# train autoregression
model = AutoReg(train, lags=29)
model_fit = model.fit()
print('Coefficients: %s' % model_fit.params)
# make predictions
predictions = model_fit.predict(start=len(train), end=len(train)+len(test)-1, dynamic=False)
for i in range(len(predictions)):
    print('predicted=%f, expected=%f' % (predictions[i], test[i]))
rmse = sqrt(mean_squared_error(test, predictions))
print('Test RMSE: %.3f' % rmse)
# plot results
plt.plot(test)
plt.plot(predictions, color='red')
plt.title("Fig. 3 Original Vs Predicted Value")
plt.xlabel("Iteration")
plt.ylabel("Value")
plt.show()

```

C:\Users\kiran\anaconda3\lib\site-packages\statsmodels\tsa\ar_model.py:248: FutureWarning: The parameter names will change after 0.12 is released. Set old_names to False to use the new names now. Set old_names to True to use the old names.

```

warnings.warn(
Coefficients: [ 2.20424835e+03  8.78872100e-01 -3.60867752e-01  3.44408784e-01
 2.72624784e-01 -2.21552823e-01  1.65049627e-01  8.35543692e-03
-1.94172346e-01  1.68899969e-01 -1.38659003e-02  3.17476649e-02
-2.05193864e-01  3.51037306e-02  1.91318789e-01 -6.44582931e-02
 2.64535613e-01 -5.15897445e-01  3.40840732e-02  3.19259770e-01
-1.57077398e-01  4.76717928e-03  1.51880104e-01 -2.81862225e-01
-1.11082273e-02  1.60691634e-01 -1.70967353e-01  2.23886941e-01
-1.37176649e-01  7.66363949e-02]
predicted=445515.494493, expected=481627.000000
predicted=422256.678222, expected=483716.000000
predicted=446066.473177, expected=493327.000000
predicted=451949.982676, expected=493991.000000
predicted=442227.695371, expected=488652.000000

```

predicted=450286.767245, expected=484782.000000
predicted=446613.649392, expected=520162.000000
predicted=438267.782201, expected=504458.000000
predicted=444921.157325, expected=559871.000000
predicted=470506.850108, expected=562269.000000
predicted=455556.374313, expected=548987.000000
predicted=424719.228433, expected=550782.000000
Test RMSE: 75538.510

Fig. 3 Original Vs Predicted Value

